



Switching and Routing

Packet classification



Guido Maier

Dipartimento di Elettronica, Informazione e Bioingegneria

Politecnico di Milano

Ph: +39 022399 3575, Fax +39 022399 3413

guido.maier@polimi.it

(*) – Acknowledgement: the basic material was provided by Jonathan Chao



Introduction

- To meet various QoS requirements, routers need to implement the following features
 - ▶ Admission control
 - ▶ Resource reservation
 - ▶ Per-flow queueing
 - ▶ Fair scheduling
- They need to be able to distinguish and classify the incoming traffic into different flows (flow-aware routers)
 - ▶ Flow-aware routing and packet classification is the founding principle of Software Defined Networking (SDN) and OpenFlow



Introduction

- Flows are specified by rules
 - ▶ Each rule consists of operations comparing packet fields with certain values
- A set of rules is called a classifier
 - ▶ Based on the criteria to be applied to classify packets with respect to a given network application




A Packet Classifier Example

Rule	IPd	IPs	Prot.	Port#	Appl	Action
R1	128.238/16	*	TCP	telnet	*	Deny
R2	176.110/16	196.27.43/24	UDP	*	RTP	Send to port III
R3	196.27.43/24	134.65/16	TCP	*	*	Drop if rate > 10 Mb/s

R1: Packet Filtering

R2: Policy Routing

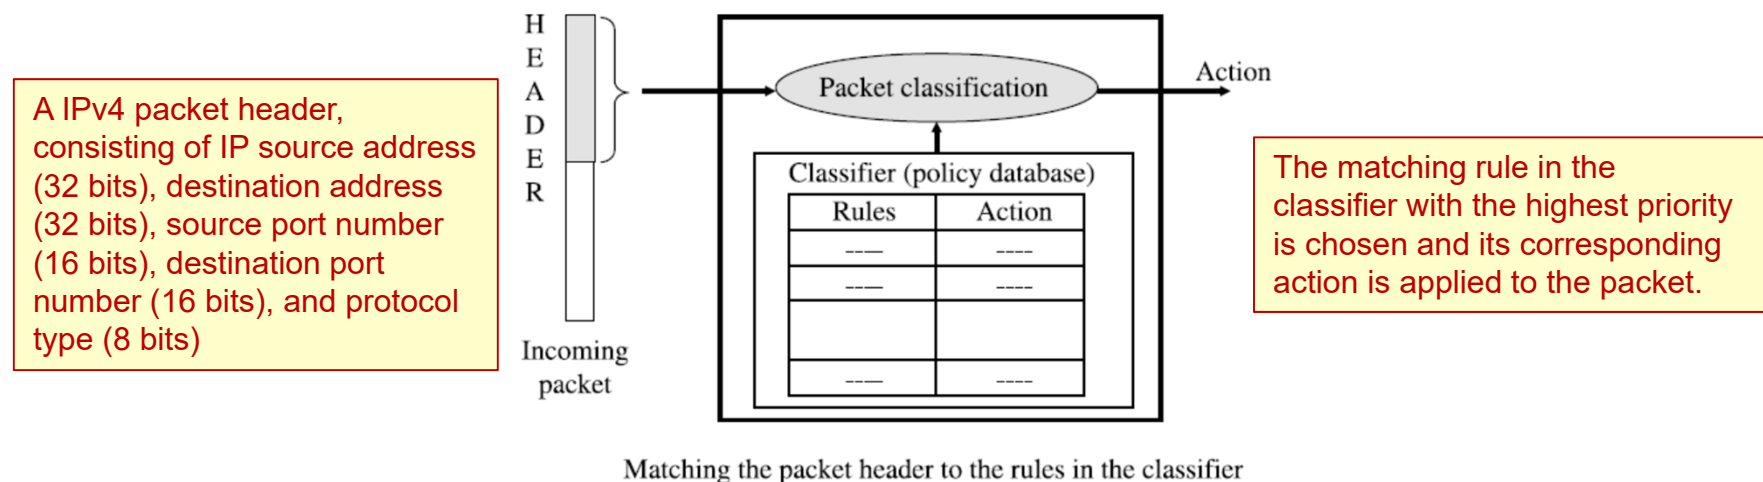
R3: Traffic Policing





Classifier definition

- A classifier C consists of N rules, R_j , $1 \leq j \leq N$, where R_j is composed of three entities:
 - ▶ (a) A regular expression $R_j[i]$, $1 \leq i \leq d$, on each of the d header fields of a packet.
 - ▶ (b) A number, $Pri(R_j)$, indicating the priority of the rule in the classifier.
 - ▶ (c) An action, referred to as $Action(R_j)$.
- An incoming packet P with the header considered as a d -tuple (P_1, P_2, \dots, P_d) is said to match R_j , if and only if, P_i matches $R_j[i]$, where $1 \leq i \leq d$.
- Given an incoming packet P and thus the d -tuple, the d -dimensional packet classification problem is to find the rule R_m with the highest priority among all the rules R_j matching the d -tuple





Classifier example

TABLE 3.1 Classifier Example

Rule	Network-Layer		Transport-Layer		Application-Layer	Action
	Destination	Source	Protocol	Destination	Protocol	
R_1	128.238/16	*	TCP	= telnet	*	Deny
R_2	176.110/16	196.27.43/24	UDP	*	RTP	Send to port III
R_3	196.27.43/24	134.65/16	TCP	*	*	Drop traffic if rate > 10 Mbps
R_4	*	*	*	*	*	Permit

- Each rule has five regular expressions on five packet-header fields from network layer to application layer
- Each expression could be
 - ▶ simple prefix/length specification → same definition as in IP lookups
 - ▶ operator/number specification → could be more general, such as equal 23, range 256-1023, and greater than 1023
- A wildcard is allowed to be inserted to match any value.
- Note R_4 'all-wildcards' specification → matches with any incoming packet
 - ▶ The priorities of rules take effect when a packet matches both R_4 and the other rules



Classifier example

TABLE 3.2 Example Classifier with Seven Rules in Four Fields

Rule	F_1	F_2	F_3	F_4	Action
R_1	00*	110*	6	(10, 12)	Act_0
R_2	00*	11*	(4, 8)	15	Act_1
R_3	10*	1*	7	9	Act_2
R_4	0*	01*	10	(10, 12)	Act_1
R_5	0*	10*	(4, 8)	15	Act_0
R_6	0*	1*	10	(10, 12)	Act_3
R_7	*	00*	7	15	Act_1

- Rule-set $C = R_j (1 \leq j \leq N)$ | each rule R_j has d fields \rightarrow fields are labeled as F_i ($1 \leq i \leq d$) and R_j is denoted as $\langle R_{j1}, R_{j2}, \dots, R_{jd} \rangle$
- Example classifier:
 - ▶ 7 rules, 4 fields, last column shows the action
 - ▶ The seven rules are listed in the order of descending priorities, that is, R_1 has the highest priority.
- F_1 and F_2 specified in prefixes \rightarrow handled more efficiently by using tries or TCAM
- F_3 and F_4 specified in ranges \rightarrow handled more efficiently by projecting the numbers into different ranges and then performing range lookup



Performance Metrics

- Several performance metrics are used to compare and analyze packet classification algorithms
- Search speed
 - ▶ In high speed links: 40-byte IP packets @ 10 Gbit/s \rightarrow 31.25 Mpack/s \rightarrow classification time < 32 ns
- Storage requirement
 - ▶ Small storage memory (cache or SRAM) \rightarrow low access time, low power consumption
- Scalability in classifier size
 - ▶ # (micro)flows in metro/edge routers: 128k - 1 M
- Scalability in the number of header fields
 - ▶ More complex services \rightarrow more header fields
- Update time
 - ▶ Some applications (e.g. flow recognition) require fast rule updating time
- Flexibility in specification
 - ▶ Wide range of rules (btw. in OpenFlow this range is very constrained)



Packet classification algorithms

- Linear search
 - ▶ The simplest algorithm for packet classification
 - ▶ Given an incoming packet header, the rules are examined one by one until a match is found
 - ▶ For a N-rule classifier, both the storage and query time complexity are $O(N)$, making this scheme infeasible for large rule sets
- Many efficient packet classification schemes have been proposed



Packet classification algorithms

- Trie-based classification
 - ▶ Hierarchical trie
 - ▶ Set-pruning trie
 - ▶ Grid of tries
 - ▶ Extending two-dimensional schemes
 - ▶ Field-level trie classification
- Geometric algorithms
 - ▶ Cross-producting scheme
 - ▶ Bitmap intersection
 - ▶ Parallel packet classification
 - ▶ Area-based quadtree
 - ▶ Hierarchical intelligent cuttings
- Heuristic algorithms
 - ▶ Recursive flow classification
 - ▶ Tuple space search
- TCAM-based algorithms



Packet classification (reduced set)

- Trie-based classification
 - ▶ Hierarchical trie
 - ▶ Set-pruning trie
 - ▶ Grid of tries
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- Heuristic algorithms
 - ▶ Recursive flow classification
- TCAM-based algorithms



Packet classification (further reduced set)

- Trie-based classification
 - ▶ Hierarchical trie
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- TCAM-based algorithms



Packet classification

- Trie-based classification
 - ▶ Hierarchical trie
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- TCAM-based algorithms



The hierarchical trie data structure


- Extension of the one-dimension trie to multiple dimension
 - ▶ Each dimension represents a field
 - ▶ Also called multi-level tries, backtracking search tries, or trie-of-tries
- Prefix-based rules can be easily processed by using tries



The hierarchical trie data structure

- Formal construction

- ▶ A binary radix trie, called F_1 -trie is first built for the set of prefixes $\{R_{j1}\}$ that belong to F_1 of all the rules
- ▶ Secondly, for each prefix p_j in the F_1 -trie, a $(d - 1)$ -dimensional hierarchical trie T_p is recursively constructed for those rules that exactly specify p in F_1 , that is, the set of rules $\{R_j \mid R_{j1} = p\}$
- ▶ Trie T_p is connected to p by a next-trie pointer stored in node p






The hierarchical trie data structure

- Classification Scheme

- ▶ Incoming packet with the header (v_1, v_2, \dots, v_d)
- ▶ Query algorithm traverses the F_1 -trie based on v_1
- ▶ If a next-trie pointer is encountered, the algorithm goes on with the pointer and queries the $(d - 1)$ -dimensional hierarchical trie recursively






The hierarchical trie data structure

C

Rule	F_1	F_2
R_1	00^*	11^*
R_2	00^*	1^*
R_3	10^*	1^*
R_4	0^*	01^*
R_5	0^*	10^*
R_6	0^*	1^*
R_7	$*$	00^*




- Nodes: Elipses $\rightarrow F_1$
-



The hierarchical trie data structure

C		
Rule	F_1	F_2
R_1	00*	11*
R_2	00*	1*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



- Nodes: Ellipses $\rightarrow F_1$; circles $\rightarrow F_2$
- 4 F_2 -tries because we have four distinct prefixes in the F_1 field of \mathcal{C}




The hierarchical trie data structure

Incoming Packet: (001, 110) F1 F2

C

Rule	F_1	F_2
R_1	00*	11*
R_2	00*	1*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



- Search process backtracks to node 'B', which is the lowest ancestor of node 'D' in the F_1 -trie
- Procedure is repeated until no ancestor node of node 'D' is available to be searched



The hierarchical trie data structure

- The backtracking process is necessary, because
 - ▶ “001” of the incoming packet may match several prefixes in the first field
 - ▶ We have no knowledge in advance which F2-trie contains prefix(es) that match “110”
- During this traversal, three matches are found, R_1 , R_2 , and R_6 . R_1 is returned as the highest priority rule matched




The hierarchical trie data structure

Incoming Packet: (010,001) ^{F1 F2}

C

Rule	F_1	F_2
R_1	00*	11*
R_2	00*	1*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



- R_7 returned



The hierarchical trie data structure

Performance


- ▶ N-rule set, each of which is with d sub-fields and the maximum field length of each field is W
- ▶ W : depth of F_d trie
- Storage complexity: $O(NdW)$
 - ▶ Is one of most storage-economic algorithms; the data structure is straightforward and easy to maintain ...
 - ▶ ... at the expenses of a longer searching time
- Search time complexity: $O(W^d)$
 - ▶ F_d -trie has a depth of W and thus takes $O(W)$ to search.
 - ▶ F_{d-1} -trie also has a depth of W , where each node has a F_d -trie.
 - ▶ The worst-case search time for the F_{d-1} -trie is thus $O(W^2)$.
 - ▶ With induction, the time complexity becomes $O(W^d)$.
- Update complexity: $O(d^2W)$
 - ▶ each field of the updated rule is stored in exactly one location in a d -level tree with maximum depth $O(dW)$

In the book is wrong!

In the book is unclear!



The set-pruning trie data structure




- Each trie node (with a valid prefix) duplicates all rules in the rule sets of its ancestors into its own rule set

Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



The set-pruning trie data structure




- Each trie node (with a valid prefix) duplicates all rules in the rule sets of its ancestors into its own rule set
- Then constructs the next dimension trie based on the new rule set

Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



The set-pruning trie data structure


- Search process for a d-tuple consists of d consecutive longest prefix matching on each dimension of the set-pruning trie
- Classification examples:
 - $(001, 110) \rightarrow$
 R_1 is returned as the highest priority rule matched
 - Multiple rules may be encountered along the path (R_6, R_2) and the one with the highest priority is recorded
 - $(001, 011) \rightarrow$
 R_4 is returned as the highest priority rule matched



Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



The set-pruning trie data structure



Rule	F_1	F_2
R_1	00*	11*
R_2	00*	1*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*

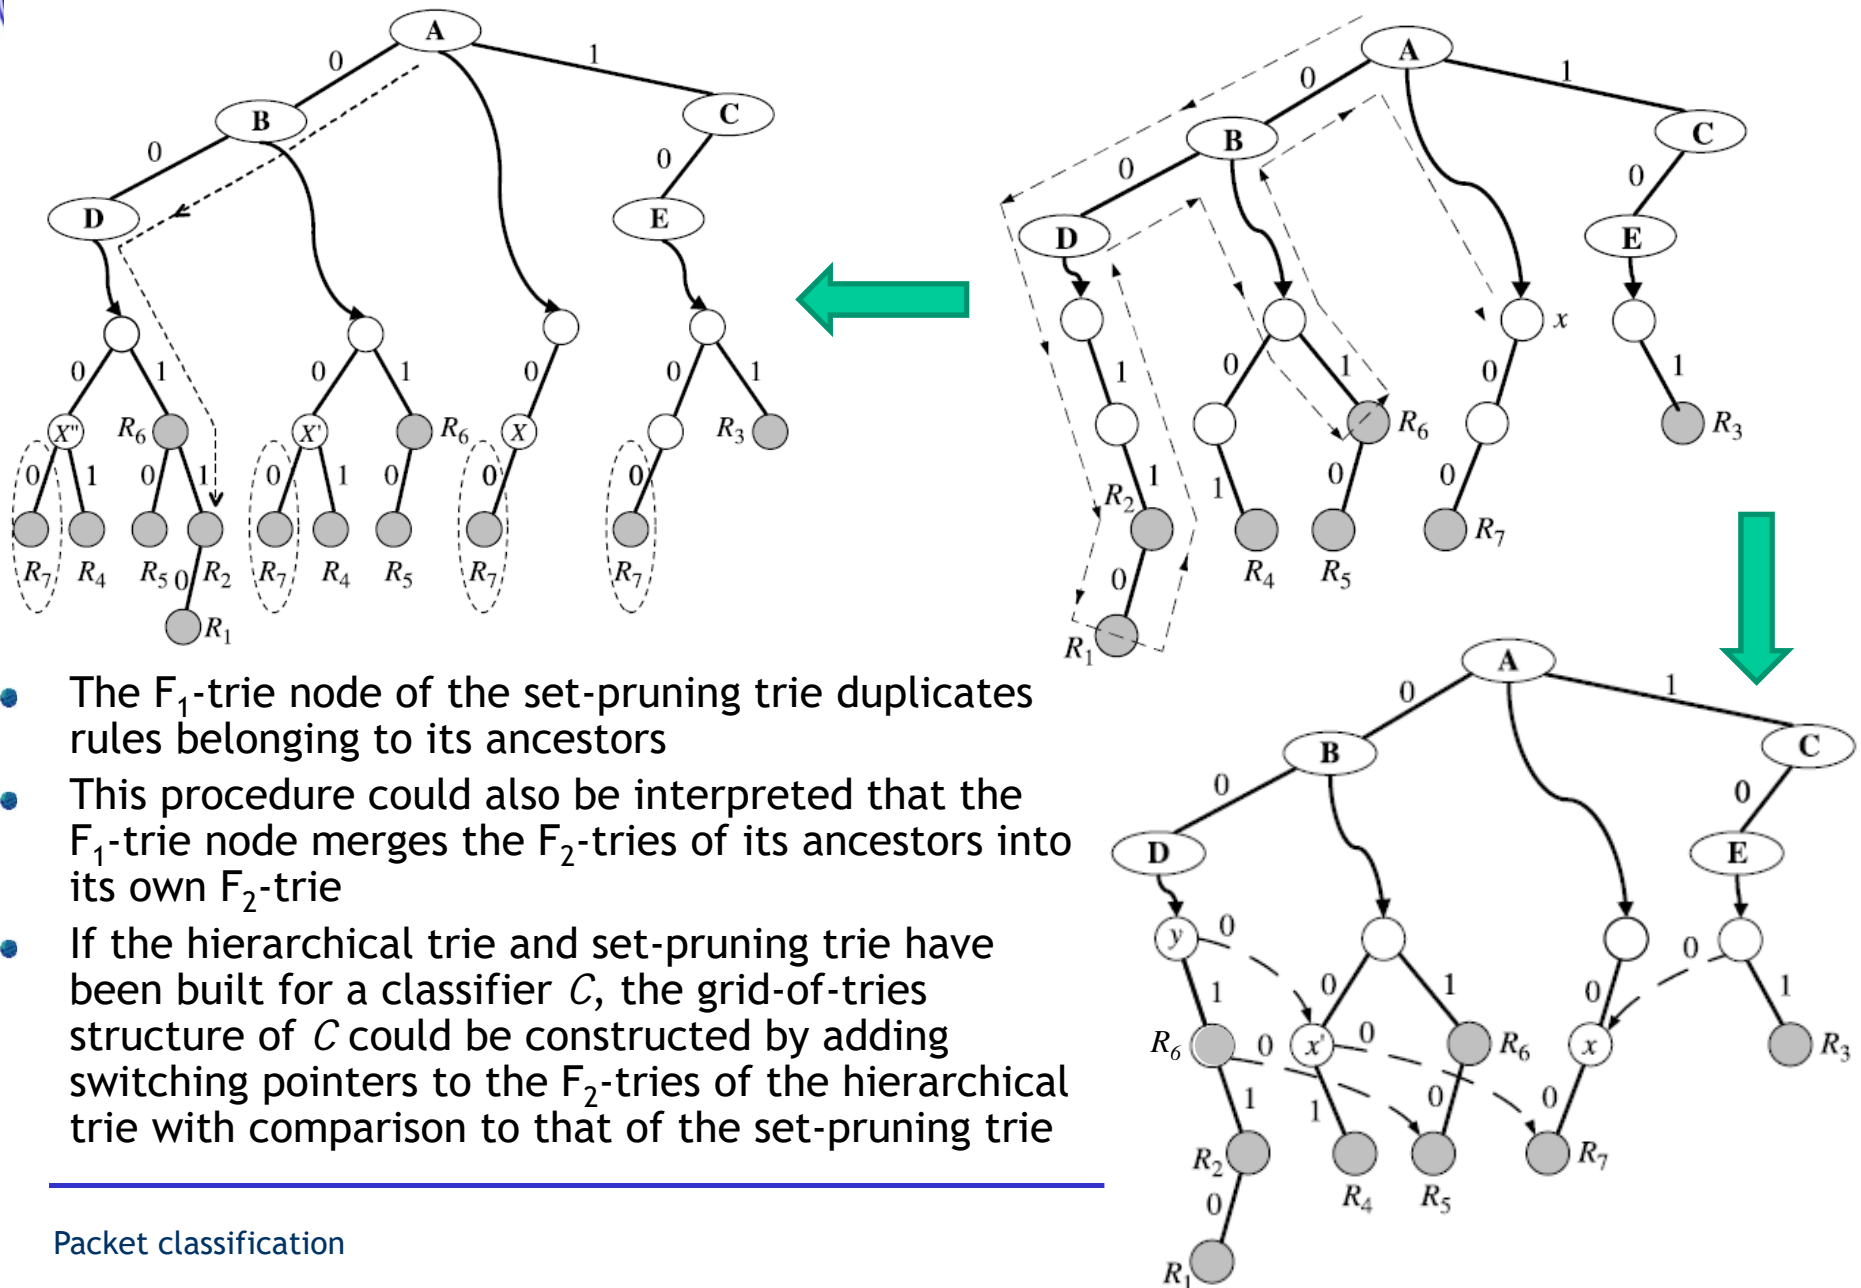
Note:
When Parent copy to Child, keep the higher priority node and drop the lower one. Here because R_2 has higher priority, drop R_6 .



The set-pruning trie data structure

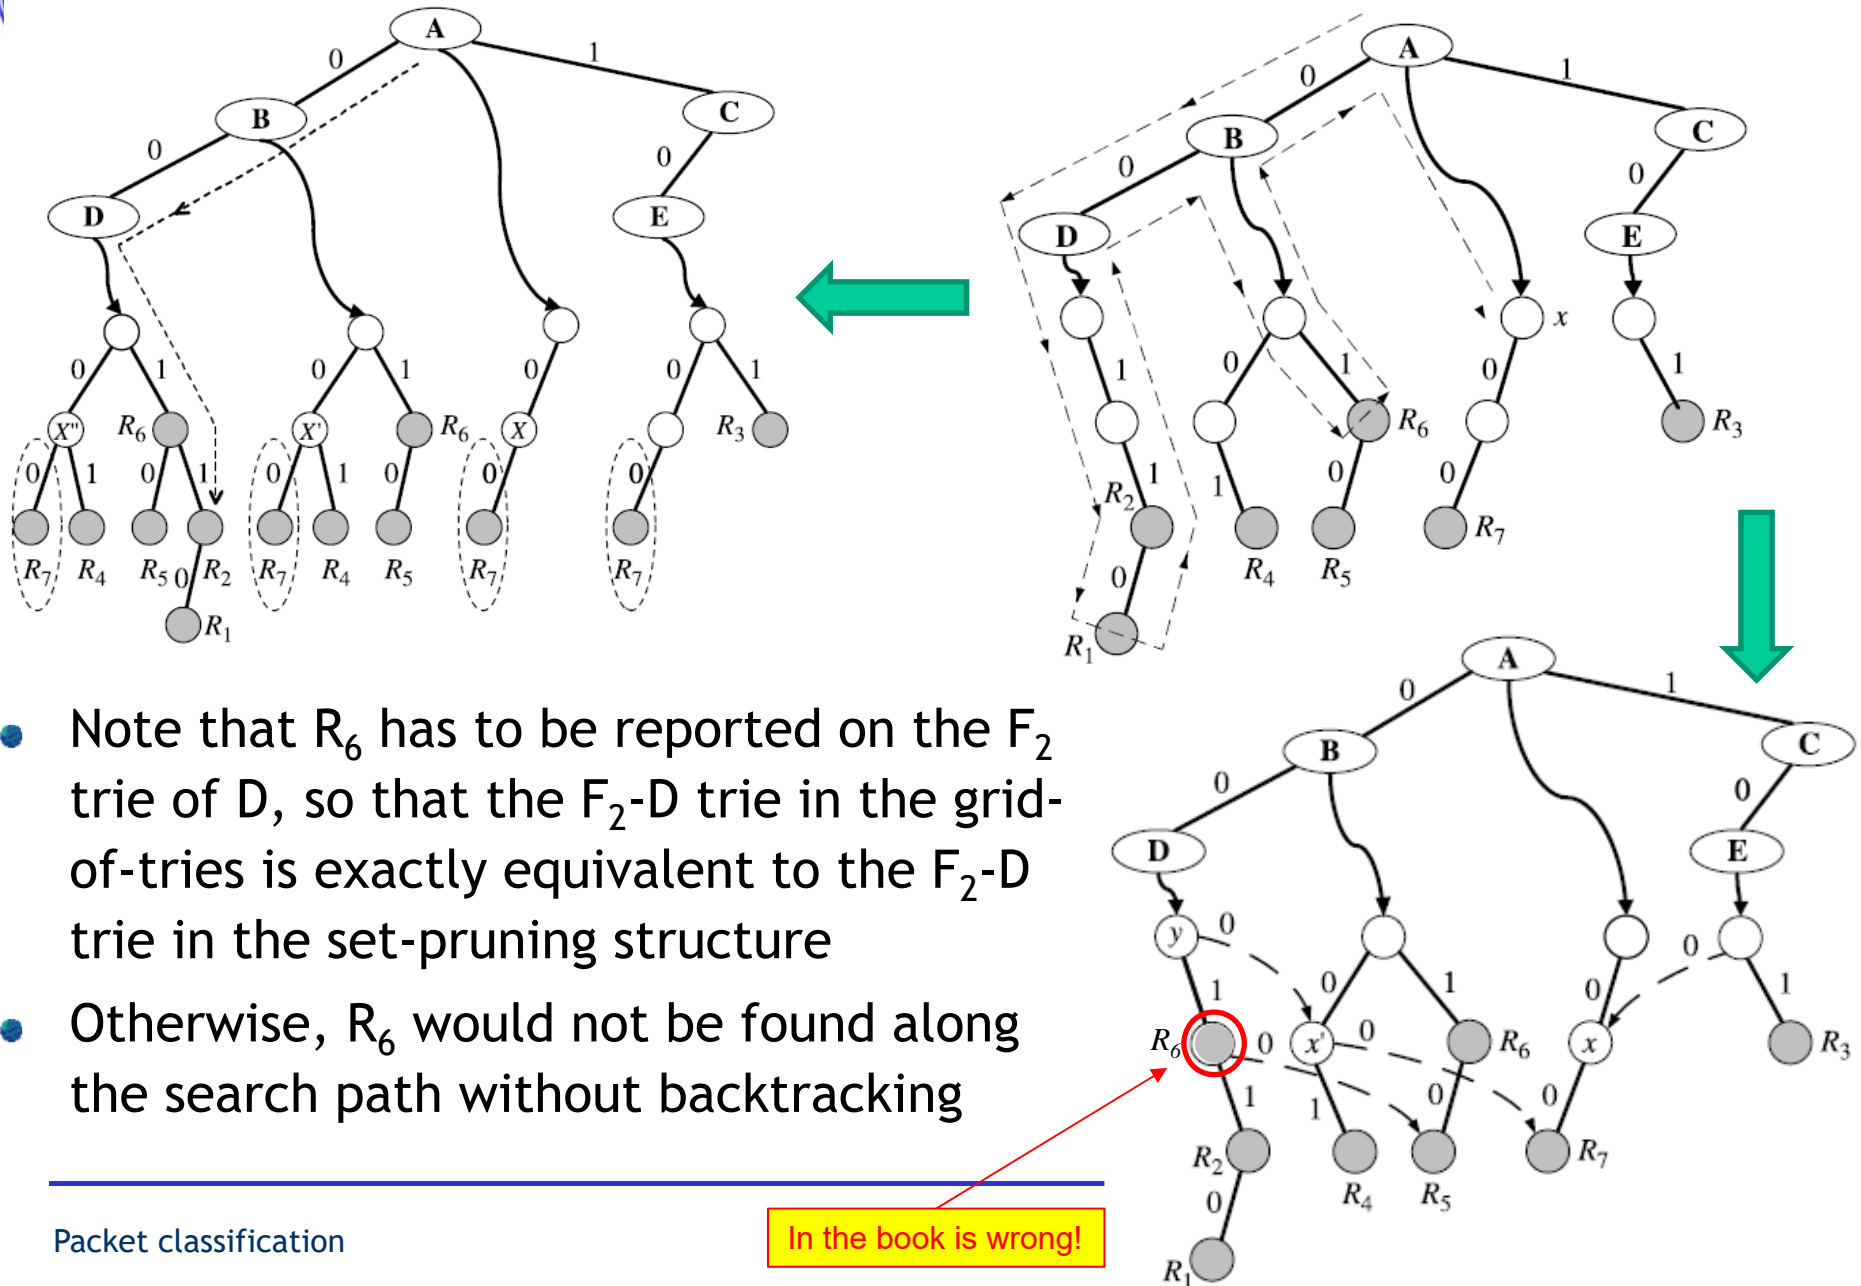
Performance

- ▶ N-rule set, each of which is with d sub-fields and the maximum field length of each field is W
- ▶ W : depth of F_d trie
- Storage complexity: $O(N^d d W)$
 - ▶ Increased by a factor N^{d-1} due to worst-case duplication of rules (a rule may be replicated N^d times)
- Search time complexity: $O(dW)$
 - ▶ Elimination of backtracking → relevant improvement compared to W^d
- Update complexity: $O(N^d)$
 - ▶ Compared to $O(d^2 W)$ of hierarchical tries







The grid-of-tries structure (for 2D class.)






The grid-of-tries structure (for 2D class.)



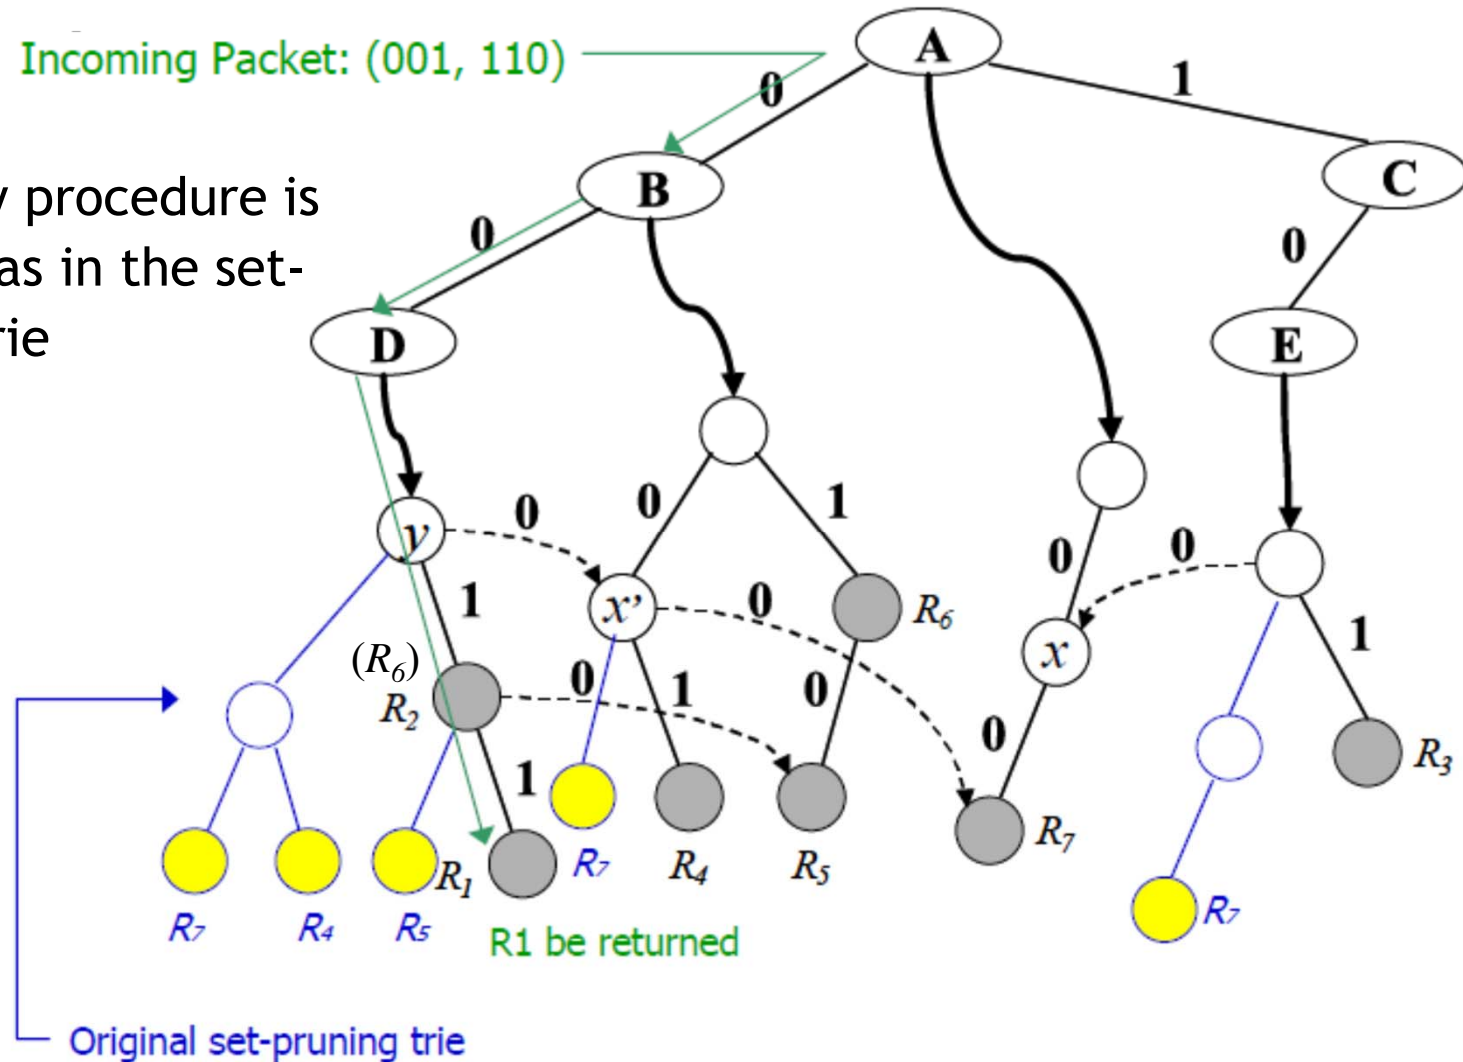
- A switching pointer, p_s , labeled with 0/1 is inserted at node y whenever its counterpart in the set-pruning trie contains a 0/1-pointer to another node z (e.g. x'') while y does not
- Node z may have several counterparts in the hierarchical trie, but p_s points to the one contained in the F_2 -trie that is 'closest' to the F_2 -trie containing node y
 - ▶ Node x and node x' are both counterparts of node x'' . However, the switching pointer at node y points to node x' since node B is closer to node D than node A





The grid-of-tries structure (for 2D class.)

- The query procedure is identical as in the set-pruning trie





The grid-of-tries structure (for 2D class.)

Performance

- ▶ W : depth of F_d trie
- ▶ N -rule set
- Storage complexity: $O(NdW) \rightarrow O(2NW)$
 - ▶ As the hierarchical trie
- Search time complexity: $O(dW) \rightarrow O(2W)$
 - ▶ As the set-pruning trie
- Update complexity:
 - ▶ Incremental updates are complex \rightarrow the entire data structure has to be rebuild, with complexity $O(NdW)$, at each update
- The grid-of-tries structure performs well on both query time and storage complexity
- But incremental updates are complex since several pointers may point to a single node
- If the node is to be removed, a new node needs to be created and the pointers need to be updated to point to the new node



Packet classification

- Trie-based classification
 - ▶ Hierarchical trie
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- TCAM-based algorithms



Geometric Algorithms

- Each field of a classifier can be specified in either a prefix/length pair or an operator/number form
- From a geometric point of view, both specifications could be interpreted by a range (or interval) on a number line
- A rule with two fields represents a rectangle in the 2D Euclidean space and a rule with d fields represents a d -dimensional hyper-rectangle
- The classifier is a set of such hyper-rectangles with priorities associated
- A packet header (d -tuple), it represents a point P in the d -dimensional space
- Packet classification problem is equivalent to finding the highest priority hyper-rectangle that encloses P

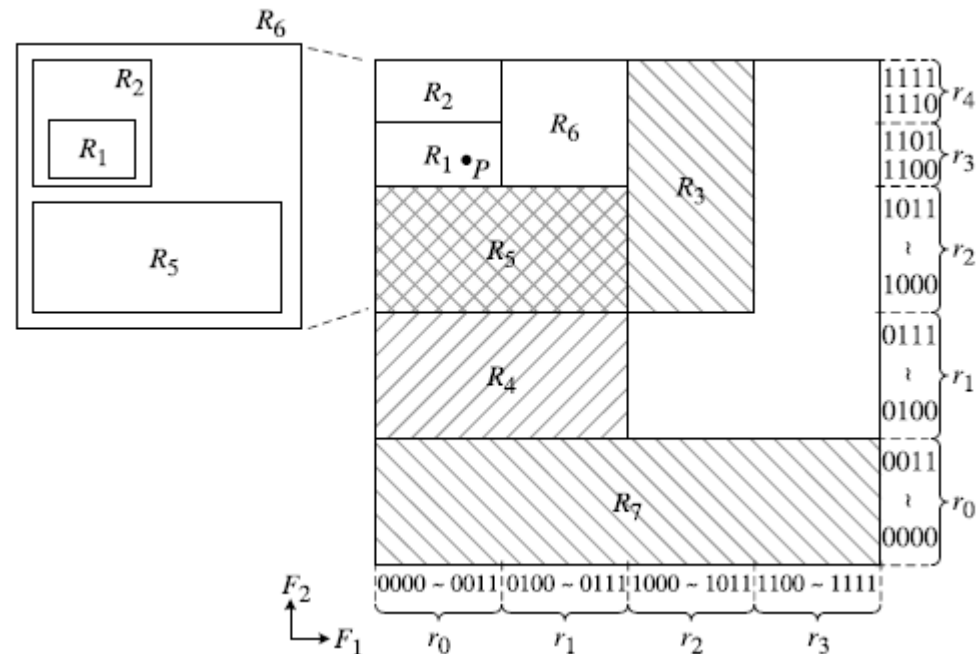


- Standard problems in the field of computational geometry that resemble packet classification
 - ▶ Point location problem: finding the enclosing region of a point, given a set of non-overlapping regions → Theoretical bounds in N (hyper-)rectangular regions and $d > 3$ dimensions are $O(\log N)$ time with $O(N^d)$ space [= memory], or $O((\log N)^{d-1})$ time with $O(N)$ space
 - ▶ Packet classification is at least as hard (regions can overlap) → implies that the packet classification is extremely complex in the worst case
 - A solution is either impracticably large (with 100 rules and 4 fields, $O(N^d)$ space is about 100MBytes) or too slow ($O((\log N)^{d-1})$ is about 290 memory accesses).



The geometric representation of the classifier

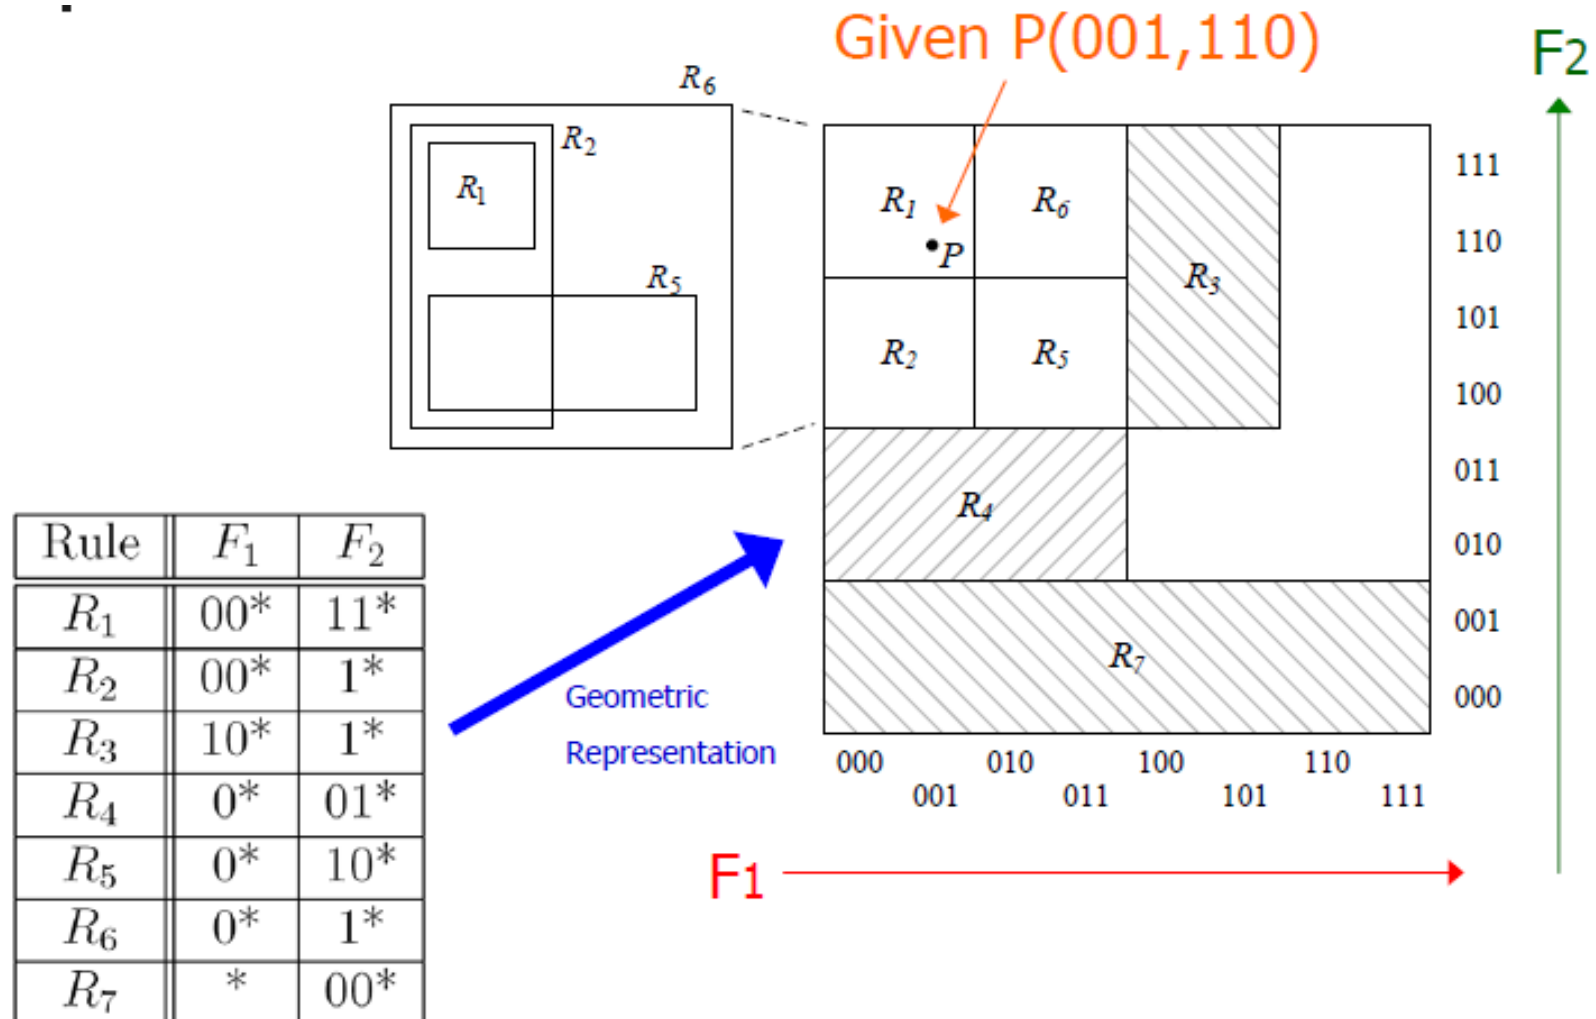
Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*



- If the prefixes or ranges in one field are projected on the number line $[0, 2^W - 1]$, a set of disjoint elementary ranges (or intervals) is obtained
 - The concatenation of these elementary ranges forms the whole number line
- Given a number Z on the number line, the range lookup problem is defined as locating the elementary range (or interval) containing Z
- A prefix represents a range on the number line



The geometric representation of the classifier (example)





The geometric representation of the cross-producting algorithm

- The cross-producting scheme works by performing d range lookup operations, one on each field
- Compose these d results to index a pre-computed table that returns the highest priority rule matched
- For instance, given a 2-tuple $(p_1; p_2)$, two range lookups are performed on each range set and the two matching ranges returned are composed to index the pre-computed table



The geometric representation of the cross-producing algorithm

Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*

	$r_1[0]$	$r_1[1]$	$r_1[2]$	$r_1[3]$	
$r_2[4]$	R_2	R_6	R_3	—	1111 1110
$r_2[3]$	R_1	R_6	R_3	—	1101 1100
$r_2[2]$	R_5	R_5	R_3	—	1011 1000
$r_2[1]$	R_4	R_4	—	—	0111 0100
$r_2[0]$	R_7	R_7	R_7	R_7	0011 0000
$F_2 \uparrow$	0000 ~ 0011 0100 ~ 0111 1000 ~ 1011 1100 ~ 1111				$F_1 \rightarrow$

- For the first step, the rule specifications in the F_1 and F_2 fields are projected on two number lines
- Two sets of ranges $\{r_1[0], \dots, r_1[3]\}$ and $\{r_2[0], \dots, r_2[4]\}$, are obtained
- Each pair of ranges $(r_1[i], r_2[j])$, corresponds to a small rectangle with a pre-computed best matching rule written inside

An array with as many entries as the number of rectangles is needed in memory



The geometric representation of the cross-producting algorithm


Rule	F_1	F_2
R_1	00*	110*
R_2	00*	11*
R_3	10*	1*
R_4	0*	01*
R_5	0*	10*
R_6	0*	1*
R_7	*	00*

	$r_1[0]$	$r_1[1]$	$r_1[2]$	$r_1[3]$	
$r_2[4]$	R_2	R_6	R_3	—	1111 1110
$r_2[3]$	R_1	R_6	R_3	—	1101 1100 ←
$r_2[2]$	R_5	R_5	R_3	—	1011 1000
$r_2[1]$	R_4	R_4	—	—	0111 0100
$r_2[0]$	R_7	R_7	R_7	R_7	0011 0000
$F_2 \uparrow$ $F_1 \rightarrow$	0000 ~ 0011	0100 ~ 0111	1000 ~ 1011	1100 ~ 1111	

- Given a 2-tuple (p_1, p_2) , two range lookups are performed on each range set and the two matching ranges returned are composed to index the pre-computed table
 - If $p_1 = 0010$ and $p_2 = 1100$, the two returned ranges $(r_1[0], r_2[3])$, tell us that R_1 is the best matching rule



The geometric representation of the cross-producting algorithm






Cross-producting algorithm performance

- The storage complexity
 - ▶ Suffers from a memory explosion problem in the worst case, when the cross-product table can have $O(N^d)$ entries
 - It can be proven that N prefixes leads to at most $2N - 2$ ranges on each dimension
- The search time complexity: $O(d t_{RL})$
 - ▶ t_{RL} is the time complexity of finding a range in one dimension.
- The update complexity
 - ▶ Incremental updates require reconstruction of the cross-product table, so it cannot support dynamic classifiers well




The geometric representation of the bitmap-insertion scheme for a 2D classifier

- The set of rules, S , that matches a packet is the intersection of d sets, S_i , where S_i is the set of rules that matches the packet in the i th dimension alone
- Four rules of a 2D classifier are depicted as four rectangles and projected on the two number lines






The geometric representation of the bitmap-insertion scheme for a 2D classifier



- Given a packet $P(p_1, p_2)$, two range lookups (e.g., using a multiway search tree) are performed in each interval set and two intervals, X_i and Y_j , which contain p_1 and p_2 , are determined



The geometric representation of the bitmap-insertion scheme for a 2D classifier



- The resulting bitmap, obtained by the intersection (a simple bitwise AND operation) of the bitmaps of X_i and Y_j , shows all matching rules for P
- If the rules are ordered in decreasing order of priority, the first `1` in the bitmap denotes the highest priority rule



The geometric representation of the bitmap-insertion scheme for a 2D classifier

- Applies to multi-dimensional packet classification with either type of specification in each field
- Based on the observation that the set of rules, S , that match a packet is the intersection of d sets, S_i , where S_i is the set of rules that match the packet in the i th dimension alone



The bitmap-insertion scheme

Performance

- The storage complexity: $O(dN^2)$
 - ▶ since each bitmap is N bits wide, and there are $O(N)$ ranges in each of the d dimensions
- The search time complexity: $O(dt_{RL} + dN/w)$
 - ▶ t_{RL} is the time to perform one range lookup
 - ▶ w is the memory width
 - ▶ Time complexity can be reduced by a factor of d by looking up each dimension independently in parallel
- The update complexity
 - ▶ Incremental updates are not well-supported
- Up to 512 rules with a 33-MHz FPGA and five 1-Mbyte SRAMs, classifying 1 mpps
 - ▶ works well for a small number of rules in multiple dimensions, but suffers from a quadratic increase in storage and linear increase in classification time with the size of the classifier



Packet classification

- Trie-based classification
 - ▶ Hierarchical trie
 - ▶ Set-pruning trie
 - ▶ Grid of tries
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- Heuristic algorithms
 - ▶ Recursive flow classification
- TCAM-based algorithms





Heuristic algorithms



- Given the complexity of the packet classification problem, a simplification is desirable by slicing the search spaces in subspaces
- The optimum is found in most of the cases, but it is not guaranteed



Recursive Flow Classification (RFC)

- Basic idea of the Recursive Flow Classification (RFC)

Classifying a packet involves mapping the H -bit packet header to a T -bit action identifier (where $T = \log N$, $T \ll H$) for a N -rule classifier.





Recursive Flow Classification (RFC)

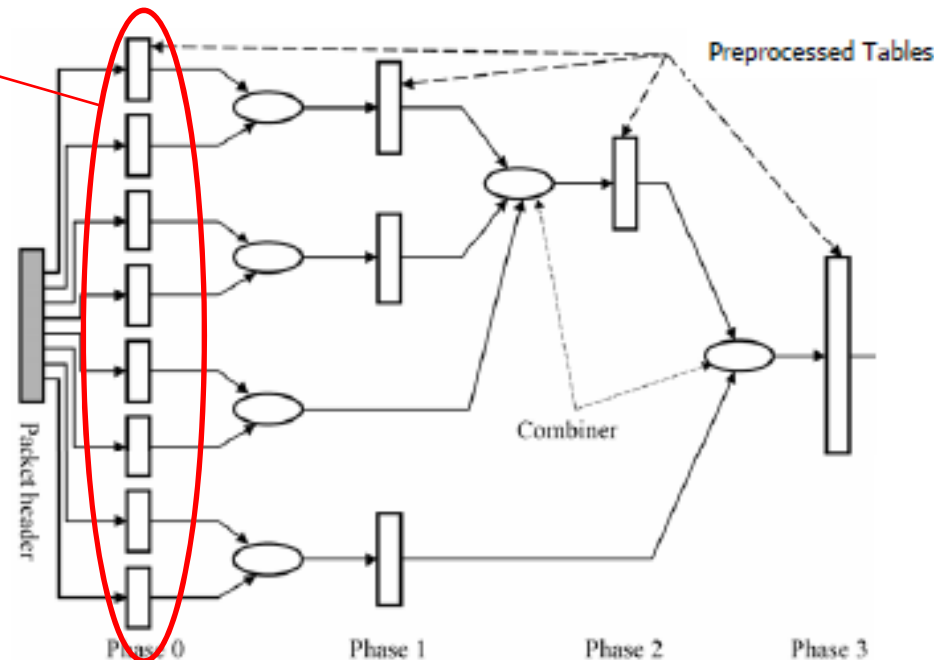
- A simple but impractical method is to pre-compute the action for each of the 2^H different packet headers and storing the results in an array $2^H \times 2^T$
 - ▶ Only one memory access is needed to yield the corresponding action. But this would require too much memory
- Main objective
 - ▶ Perform the same mapping but over several stages (phases)



Recursive Flow Classification (RFC)

- In the first step the d fields of the packet header are split into multiple chunks that are used to index multiple memories in parallel
- The mapping is performed recursively; at each stage, the algorithm performs a reduction, mapping one set of values to a smaller set
- In each phase, a set of memories returns a value shorter expressed in fewer bits, than the index of the memory access
- Results of a phase are combined to be used as input to the subsequent phase

Number of
chunks = 8





Recursive Flow Classification (RFC)

- Structure and Packet flow in the recursive flow classification

Destination IP (addr/mask)	Source IP (addr/mask)	Port Number	Protocol
152.163.190.69/0.0.0.0	152.163.80.11/0.0.0.0	*	*
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	eq www	UDP
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	range 20-21	UDP
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	eq www	TCP
152.163.198.4/0.0.0.0	152.163.160.0/0.0.3.255	gt 1023	TCP
152.163.198.4/0.0.0.0	152.163.36.0/0.0.0.255	gt 1023	TCP

The construction of the preprocessed tables is explained with a sample rule set C .


Preprocessed Tables

gt = greater than = >

- The first step of the preprocessed table construction
 - Split the d fields of the packet header into multiple chunks that are used to index multiple memories in parallel

Chopping of packet header into chunks for rule set R in the first phase of RFC:

	Source IP	Destination IP	Port	Protocol
Size (bit)	32	32	16	8
Chunk #	0	1	2	3






Recursive Flow Classification (RFC)

- Each of the parallel lookups will map the chunk to an *eqID* according to the rules
 - Consider a fixed chunk of size b bits. Its mapping table is of 2^b entries and each entry contains an *eqID* for that chunk value
 - The *eqIDs* are determined by those component(s) of the rules in the classifier corresponding to this chunk
 - The term "Chunk Equivalence Set" (CES) is used to denote a set of chunk values that have the same *eqID*
 - e.g. for chunk #3 (protocol, 8 bits), the 'CES' will be:
 - (a) {TCP}
 - (b) {UDP}
 - (c) {all remaining numbers in the range 0-255}

Protocol
*
UDP
UDP
TCP
TCP
TCP





Recursive Flow Classification (RFC)


- Each CES can be constructed in the following manner:
 - ▶ For a b -bit chunk, project the rules in the classifier on to the number line $[0, 2^b-1]$
 - ▶ Each component projects to a set of intervals on the number line
 - ▶ The end points of all the intervals projected by these components form a set of non-overlapping intervals
 - Two points in the same interval always belong to the same equivalence set
 - Two intervals are also in the same equivalence set if exactly the same rules project onto them



Recursive Flow Classification (RFC)


Example:

- Computing the four equivalence classes $E_0 \dots E_3$ for chunk #2 (corresponding to the 16-bit transport-layer destination port number) in the rule set



The four CESs can be decoded using two bits. For example we can assign '00' to E_1 , '01' to E_0 , '10' to E_2 and '11' to E_3 . Then the RFC table for this chunk is filled with the corresponding *eqIDs*, such as $table(20) = 01'$, $table(23) = 11'$, etc.

Port Number	l
*	
eq www	
range 20-21	
eq www	
gt 1023	
gt 1023	




Well-known port number for http (www): 80



Recursive Flow Classification (RFC)

- A chunk is formed by a combination of two (or more) chunks obtained from memory lookups in previous steps, with a corresponding CES
 - ▶ If the resulting chunk is of width b bits, we again create equivalence sets
 - Two b -bit numbers that are not distinguished by the rules of the classifier belonging to the same CES
 - ▶ Compute all possible intersections of the equivalence sets from the previous steps being combined
 - Each distinct intersection is an equivalence set for the newly created chunk





Recursive Flow Classification (RFC)


Example

- ▶ If we combine chunk #2 (port number) and #3 (protocol), then five CES scan be get:

- $\{(\{80\}, \{\text{UDP}\})\}$
- $\{(\{20-21\}, \{\text{UDP}\})\}$
- $\{(\{80\}, \{\text{TCP}\})\}$
- $\{(\{\text{gt}1023\}, \{\text{TCP}\})\}$
- {all the remaining crossproducts}

Destination IP (addr/mask)	Source IP (addr/mask)	Port Number	Protocol
152.163.190.69/0.0.0.0	152.163.80.11/0.0.0.0	*	*
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	eq www	UDP
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	range 20-21	UDP
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0	eq www	TCP
152.163.198.4/0.0.0.0	152.163.160.0/0.0.3.255	gt 1023	TCP
152.163.198.4/0.0.0.0	152.163.36.0/0.0.0.255	gt 1023	TCP

- ▶ These can be expressed in a three-bit eqID, as shown in following Figure



CES Port	CES Prot	Port Number and Protocol	Class Number	eqID (only 3 bits required)
01	01	eq www & udp	1	000
00	01	Range 20-21 & udp	2	001
01	00	eq www & tcp	3	010
10	00	gt 1023 & tcp	4	011
....	all remaining crossproducts	5	100

- ▶ We can see that, during step two the number of bits has been reduced from four (two bits for chunk #2 and #3 respectively after step one) to three
- ▶ For the combination of the two steps, this number has dropped from 24 to 3



Recursive Flow Classification (RFC)

Port Number	Class Number	eqID (only 2 bits required)
Range 20-21	1	00
eq www	2	01
gt 1023	3	10
0-19,22-79,81-1023	4	11

(c) Port number field made into chunks and eqIDs

Protocol	Class Number	eqID (only 2 bits required)
tcp	1	00
udp	2	01
all remaining protocols	3	10

(d) Protocol field made into chunks and eqIDs


CES Port	CES Prot	Port Number and Protocol	Class Number	eqID (only 3 bits required)
01	01	eq www & udp	1	000
00	01	Range 20-21 & udp	2	001
01	00	eq www & tcp	3	010
10	00	gt 1023 & tcp	4	011
....	all remaining crossproducts	5	100

(e) Port number and protocol fields combined and made into chunks and eqIDs



Recursive Flow Classification (RFC)

- Rule storing organization for RFC for the rule set in Table



Destination IP (addr/mask)	Source IP (addr/mask)
152.163.190.69/0.0.0.0	152.163.80.11/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.163.198.4/0.0.0.0	152.163.160.0/0.0.3.255
152.163.198.4/0.0.0.0	152.163.36.0/0.0.0.255

Note that the entries are not prefixes

Destination IP (addr/mask)	Class Number	eqID (only 2 bits required)
152.163.190.69/0.0.0.0	1	00
152.168.3.0/0.0.0.255	2	01
152.163.198.4/0.0.0.0	3	10

(a) Destination IP field made into chunks and eqIDs


Source IP (addr/mask)	Class Number	eqID (only 2 bits required)
152.163.80.11/0.0.0.0	1	00
152.163.200.157/0.0.0.0	2	01
152.163.160.0/0.0.3.255	3	10
152.163.36.0/0.0.0.255	4	11

(b) Source IP field made into chunks and eqIDs



Recursive Flow Classification (RFC)

- Rule storing organization for RFC for the rule set in Table



Destination IP (addr/mask)	Source IP (addr/mask)
152.163.190.69/0.0.0.0	152.163.80.11/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.168.3.0/0.0.0.255	152.163.200.157/0.0.0.0
152.163.198.4/0.0.0.0	152.163.160.0/0.0.3.255
152.163.198.4/0.0.0.0	152.163.36.0/0.0.0.255

CES D-IP	CES S-IP	CES S-IP + D-IP (3 bit)
00	00	000
01	01	001
10	10	010
10	11	011
....	100



Recursive Flow Classification (RFC)

Classification scheme

- First split into several chunks to be used as an index
- Then the required *eqIDs* are combined into chunks of the second phase
- This procedure goes on until the final phase is reached
 - ▶ When all the remaining *eqIDs* have been combined into only one chunk
- The corresponding table will hold the actions for that packet



Recursive Flow Classification (RFC)

Performance

- The storage complexity
 - ▶ Different combination of the chunks can yield different storage requirement
 - ▶ With real-life 4D classifiers of up to 1700 rules
 - RFC appears practical for 10-Gbps line rates in hardware and 2.5-Gbps rates in software
 - The storage space and preprocessing time grow rapidly for classifiers larger than 6000 rules
 - ▶ An optimization reduces the storage requirement of a 15,000 four-field classifier to below 4 Mbytes



Packet classification

- Trie-based classification
 - ▶ Hierarchical trie
- Geometric algorithms
 - ▶ Cross-producing scheme
 - ▶ Bitmap intersection
- TCAM-based algorithms



TCAM-Based Algorithms

- Ternary content addressable memory (TCAM) → increasing popularity for fast packet classification
- TCAM coprocessor works as a lookaside processor on behalf of a network processor
 - ▶ The processor generates a search key based on the information from the packet header
 - ▶ Passes it to the TCAM coprocessor for classification via a NPU/TCAM coprocessor interface
 - ▶ Each entry of the TCAM stores a rule by concatenating the prefixes of all fields
 - ▶ TCAM coprocessor finds a matched rule in $O(1)$ clock cycles
 - Highest possible lookup/matching performance
- A local CPU is in charge of rule table update through a separate CPU-TCAM coprocessor interface

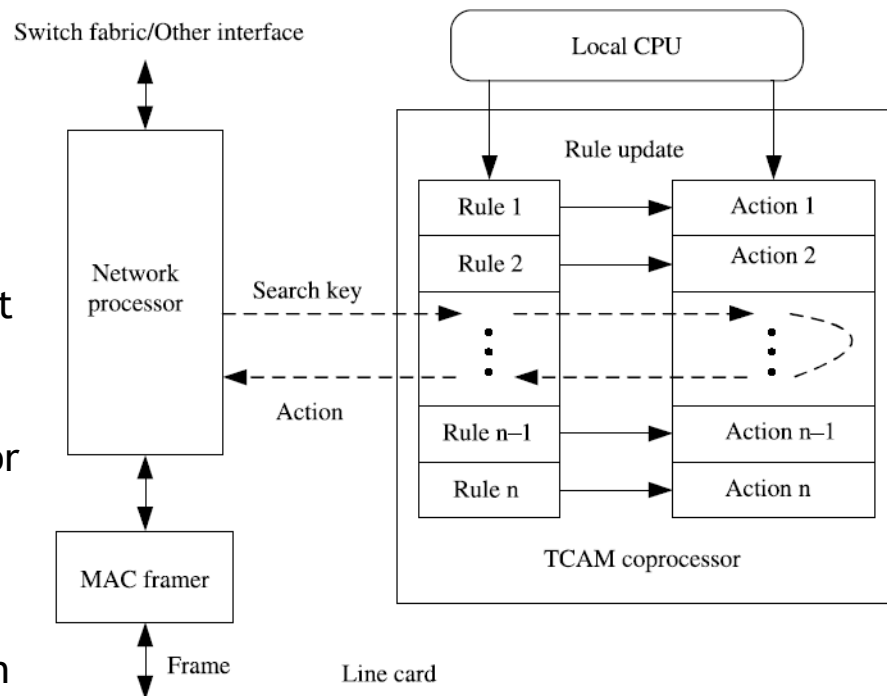


Figure 3.25 Network processor and its TCAM coprocessor.



TCAM-Based Algorithms

Range splitting

- TCAMs have been widely used in IP route lookup for longest prefix matching,
 - ▶ Rules are stored as $(val, mask)$ pairs in decreasing order of prefix lengths
- For range matching, utilization is more critical
- Range splitting must be performed to convert the ranges into prefix formats to fit the bit boundary
 - ▶ Increases the number of entries by a factor of $O(W^d)$ in the worst case
 - ▶ A range is said to be exactly implemented in a TCAM if it is expressed in a TCAM without being encoded
 - E.g. six rule entries are needed to express a range
 $\{>1023\} = \{> 0000\ 0011\ 1111\ 1111\}$




Figure 3.26 Range >1023 is expressed in terms of 6 sub-ranges in a TCAM.



TCAM-Based Algorithms

Rule / key encoding

- Reduced TCAM memory efficiency due to range matching (up to even 16%) → power consumption, footprint, and cost serious concern
- Solution:
Range preprocessing/encoding by mapping ranges to a short sequence of encoded bits, known as bit-mapping
- Rule encoding (pre-processed in software)
 - ▶ view a d-tuple rule as a region in a d-dimensional rule space
 - ▶ encode any distinct overlapped regions among all the rules
 - ▶ each rule can be translated into a sequence of encoded bits
- Search key encoding (performed on a per packet basis in hardware at wire-speed)
 - ▶ the information is extracted from the packet header
 - ▶ a search key is preprocessed to be encoded
- The encoded search key is matched against all the encoded rules to find the best matched rule



TCAM-Based Algorithms

Rule encoding for range mapping

- An efficient mapping scheme of range classifier into TCAM expands TCAM horizontally (using more bits per entry)
- For width limited application, another algorithm allows both horizontal and vertical expansion
- Rule storing organization: For each range field, an n bits vector $B = \{b_1, b_2, \dots, b_n\}$ is used to represent it
 - ▶ n is the number of distinct ranges specified for this field
 - ▶ The B vector for a range E_i has 1 at bit position i , i.e., $b_i = 1$ and all other bits are set to don't care
 - The number of distinct ranges specified for any range field is very limited
 - Exact match specification ($n = 1$) also happens frequently for a range field



TCAM-Based Algorithms

Rule encoding

- Example

R_i	Dest IP Addr (IP/mask)	Dest Port Range	Action
1	10.0.0.0/255.0.0.0	> 1023	Deny
2	192.168.0.0/255.255.0.0	50–2000	Allow
3	192.169.0.0/255.255.0.0	80(http)	High Priority
4	172.16.0.0/255.255.0.0	23(telnet)	Route through Port A
5	172.16.0.0/255.255.0.0	21(ftp)	Rate Limit to 1Mbit/s

The range of greater than 1023 in R_1 is represented By 'xxxx1'.



The bit vector representation ($n=5$ in this case).

R_i	TCAM rules	
1	10.x.x.x	xxxx1
2	192.168.x.x	xxx1x
3	192.169.x.x	xx1xx
4	172.16.x.x	x1xxx
5	172.16.x.x	1xxxx



TCAM-Based Algorithms

Search key encoding

- Classification Scheme
 - ▶ A lookup key $v \in [0, 2^k]$ is translated into an n bit vector $V = \{v_1, v_2, \dots, v_n\}$
 - ▶ Bit v_i is set to 1 if the key v falls into the corresponding range E_i , otherwise it will be set to 0
- Lookup key translation could be implemented as a direct memory lookup since most of the range fields are less than 16-bit wide



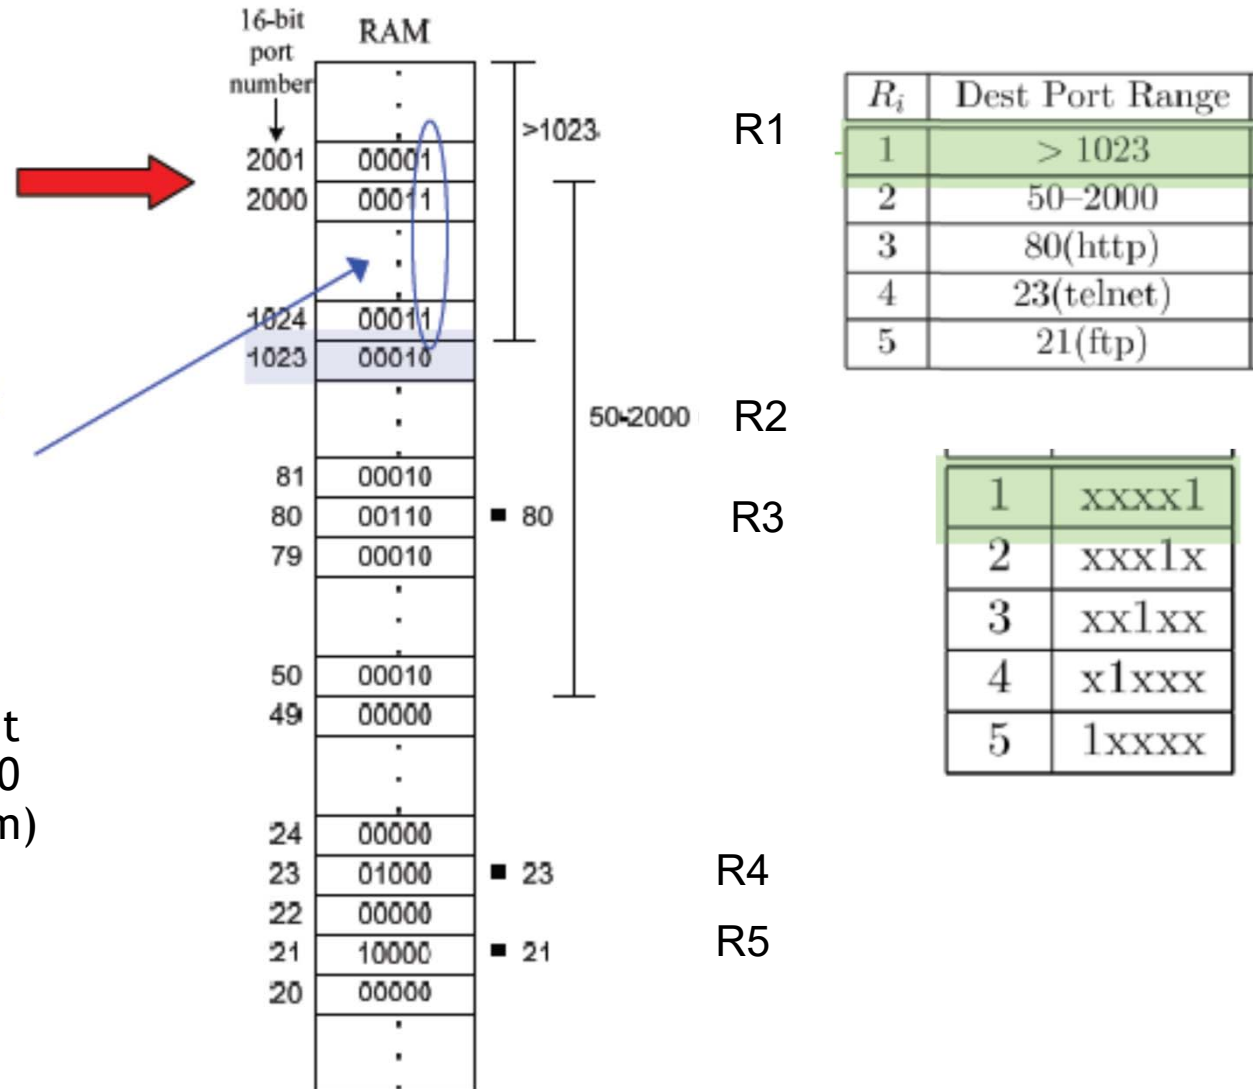
TCAM-Based Algorithms

Lookup key translation table

A complete lookup key translation table for the example classifier for each possible lookup key value.

For example, the most right of the bit vector is set to 1 for all the locations above 1023.

Second bit from the right between the locations 50 and 2000 (including them) is also set to 1 ... etc.





TCAM-Based Algorithms

Exact matching optimization

- It is possible to reduce the number of bits used to $\log_2(m + 1)$, where m is the number of exact matches
- The bit representation will contain two parts $\langle B_e, B \rangle$: B_e for exact matches, and B for all others
 - ▶ $B_e = \{b_1, b_2, \dots, b_t\}$ is a t bit vector, where $t \geq \log_2(m + 1)$ and m is the number of exact matches
 - ▶ For a normal range, $\langle B_e=0, B \rangle$ and its B portion is the same as before
 - ▶ For an exact match, $\langle B_e=i, B=0 \rangle$, if it is the i -th exact match



TCAM-Based Algorithms

Exact matching optimization

- Example

- ▶ If we use bit 2 and 3 as B_e (where bit 1 is the left most significant bit), the resulting classifier stored in TCAM will as shown on the right side table
- ▶ only two bits are needed to represent the three distinct exact matches

R_i	TCAM rules	
1	10.x.x.x	xxxx1
2	192.168.x.x	xxx1x
3	192.169.x.x	xx1xx
4	172.16.x.x	x1xxx
5	172.16.x.x	1xxxx

R_i	TCAM rules	
1	10.x.x.x	xxxx1
2	192.168.x.x	xxx1x
3	192.169.x.x	x01xx
4	172.16.x.x	x10xx
5	172.16.x.x	x11xx





TCAM-Based Algorithms

Exact matching optimization

- The lookup key translation table needs to be changed accordingly
- It also contains two parts $\langle V_e, V \rangle$ V_e corresponding to all exact matches and V to the rest
 - ▶ $V_e = \{v_1, v_2, \dots, v_t\}$ is a t bit vector
 - ▶ $V_e = I$ if the lookup key v equals to the i -th exact match, otherwise $V_e = 0$



TCAM-Based Algorithms

Lookup example

192.169.10.1 / 80



$\langle V_e, V \rangle = \langle \{01\}, 10 \rangle$

R_i	TCAM rules	
1	10.x.x.x	xxxx1
2	192.168.x.x	xxx1x
3	192.169.x.x	x01xx
4	172.16.x.x	x10xx
5	172.16.x.x	x11xx

R3

- A new packet arrives with destination IP address 192.169.10.1 and port number 80
 - ▶ The port number is indexed into the lookup key translation table
 - ▶ The resulting $V = 10$ because 80 falls in range 50-2000 but not range >1023
 - ▶ The resulting $V_e = 01$ because 01 is the value assigned to exact match value of 80
 - ▶ Together with destination IP address, the final result is rule 3



TCAM-Based Algorithms

Performance

- The storage complexity
 - ▶ This scheme requires less TCAM storage space
 - ▶ It can accommodate a much larger number of rules in a single TCAM table and reduce system cost and power consumption
- Adding/Deleting rules cause the changes of bit vectors and may require re-computation of the entire translation table, a time consuming process